Introduction to Deep Learning

Mayank Pratap Singh

August 2, 2024

1 Types of Neural Networks

1.1 Multi-layer Perceptron (MLP)

- Primarily used in supervised learning problems.

1.2 Convolutional Neural Network (CNN)

- Utilized in image and video processing.
- Developed by Yann LeCun.

1.3 Recurrent Neural Network (RNN)

- Suitable for sequential data processing.
- Includes variations like Long Short-Term Memory (LSTM) networks.

1.4 Autoencoders

- Used for data compression without significant loss of quality.

1.5 Generative Adversarial Networks (GANs)

- Developed by Ian Goodfellow.
- Consist of generators and discriminators.
- Applications include generating music, stories, and images.

2 History of Deep Learning

2.1 Early Foundations (1940s-1960s)

- **1943**: Warren McCulloch and Walter Pitts created a computational model of neural networks based on threshold logic.
- **1958**: Frank Rosenblatt introduced the perceptron, a simple two-layer neural network for pattern recognition.
- **1960s**: The concept of backpropagation emerged, with contributions from Henry Kelley and Stuart Dreyfus.

2.2 Growth and Challenges (1970s-1980s)

- **1970s**: The first AI winter led to reduced funding and interest in AI and neural networks.
- **1980**: Kunihiko Fukushima proposed the Neocognitron, a hierarchical neural network.
- **1986**: Geoffrey Hinton demonstrated that backpropagation could effectively train multi-layer networks.

2.3 Resurgence and Breakthroughs (1990s-2010s)

- **1997**: LSTM networks were developed by Sepp Hochreiter and Jürgen Schmidhuber.
- **2006**: Hinton's work on deep belief networks showed layer-by-layer training of deep architectures.
- **2012**: A deep learning model won the ImageNet competition, achieving human-level accuracy in object recognition.

2.4 Modern Era and Applications (2010s-Present)

- **2014**: Google acquired DeepMind.
- **2016**: DeepMind's AlphaGo defeated a world champion Go player.
- **Ongoing Developments**: Deep learning is integral to computer vision, natural language processing, and speech recognition.

3 Key Developments in Deep Learning History

3.1 The 1960s: Perceptron Research

- Frank Rosenblatt's work on the perceptron.
- In 1969, Marvin Minsky showed that perceptrons couldn't solve non-linear functions like XOR, leading to an AI winter.

3.2 The 1980s: Revival and Backpropagation

- **1986**: Geoffrey Hinton co-authored a paper titled "Learning representations by back-propagating errors," popularizing the backpropagation algorithm.
- **1989**: Yann LeCun's work on handwritten digit recognition and the MNIST dataset.

3.3 2006: Unsupervised Pre-training

- Geoffrey Hinton published a paper on unsupervised pre-training, initializing models before fine-tuning them on labeled datasets.

3.4 2012: The ImageNet Challenge

- AlexNet, a deep convolutional neural network, won the ImageNet competition with a top-5 error rate of 15.3%.

3.5 Generative Adversarial Networks (GANs)

- Developed by Ian Goodfellow in 2014.
- Consist of a generator and a discriminator.

4 Applications of Deep Learning

- Self-driving cars: Enabling perception, decision-making, and navigation.
- Game-playing agents: Examples include Google's DeepMind AlphaGo.
- Virtual assistants.
- Image colorization: Converting black-and-white images to color.
- Adding audio to silent videos.
- Image generation.
- Real-time text translation.
- Pixel restoration: Converting blurry images into high-quality images.
- Object detection.

5 Factors Contributing to the Popularity of Deep Learning

5.1 Datasets

- Exponential data generation due to smartphones and affordable internet.
- Large companies label data and make it publicly available.

5.2 Hardware

- Moore's Law: Transistor sizes reduce every two years.
- Significant increase in RAM and processing power in modern devices.
- Deep Learning requires high computational power for matrix operations, which are efficiently handled by GPUs.
- Advanced hardware for deep learning includes GPUs, TPUs, FPGAs, ASICs, and NPUs.

5.3 Frameworks/Libraries

- Development of accessible frameworks/libraries has simplified the training process.
- Key frameworks:
  - TensorFlow (Google)
  - PyTorch (Facebook)
- Integration of Keras with TensorFlow has made it more accessible.
- Other tools include AutoML and Apple's CreateML for simplified model building.

5.4 Architectures

- Various deep learning architectures designed for specific tasks.
- Examples of architectures:
  - ResNet for image classification
  - BERT for text classification
  - U-Net for image segmentation
  - Pix2Pix for image-to-image translation
  - YOLO for object detection
  - WaveNet for speech generation
- Transfer learning: Using pre-trained models for specific tasks instead of training from scratch.

5.5 Community

- Deep learning has a vibrant community of researchers, practitioners, and enthusiasts.
- Contributions from the community include detailed notes, tutorials, and resources, often made using LaTeX for documentation.